

Building a Batch Analytics Pipeline on HDFS & Hive

1. Introduction:

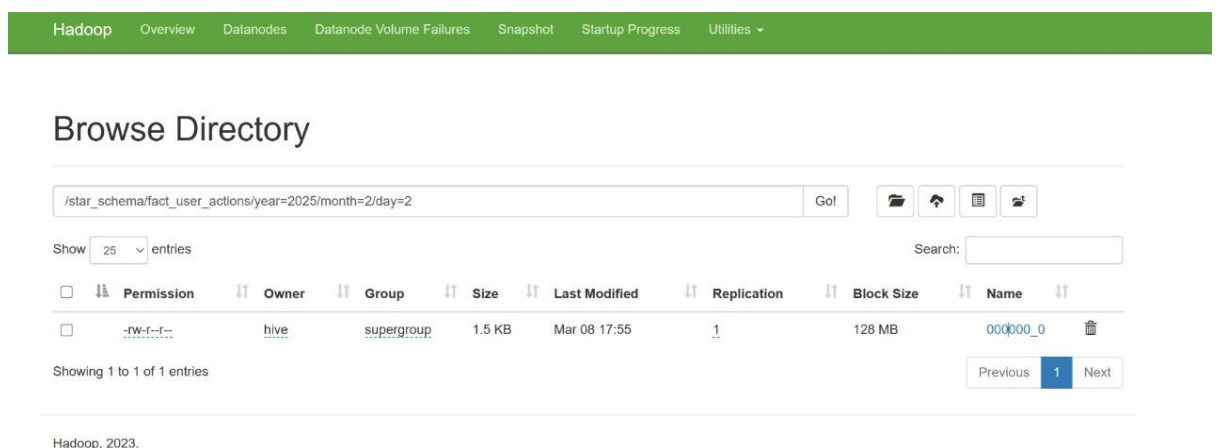
The report outlines the implementation of a data pipeline using Hadoop and hive to process user activity logs and content metadata effectively. We also follow a star schema design.

2. Data ingestion and Storage:

Raw data storage: we store the data known as raw_data which contains the input files before ingestion.

The data is ingested into HDFS under directory /raw/logs/ and /raw/metadata.

It is automated in the shell script by ingest_logs.sh



3. Hive Schemas Definitions:

Raw tables:

```
CREATE EXTERNAL TABLE IF NOT EXISTS user_activity_logs (  
  user_id STRING,  
  action STRING,  
  `timestamp` STRING,  
  details STRING  
)  
PARTITIONED BY (year INT, month INT, day INT)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION 'hdfs://localhost:9000/raw/logs';
```

Optimized Star Schema (Parquet)

```
CREATE TABLE IF NOT EXISTS dim_users (  
  user_id INT,  
  device STRING,  
  user_region STRING)  
STORED AS PARQUET;
```

```

CREATE TABLE IF NOT EXISTS dim_media (
  media_id INT,
  media_title STRING,
  genre STRING,
  duration INT,
  creator STRING)
STORED AS PARQUET;

CREATE TABLE IF NOT EXISTS dim_sessions (
  session_id STRING,
  user_id INT)
STORED AS PARQUET;

CREATE TABLE IF NOT EXISTS fact_activity_events (
  user_id INT,
  media_id INT,
  session_id STRING,
  activity_type STRING,
  activity_timestamp STRING)
PARTITIONED BY (year INT, month INT, day INT)
STORED AS PARQUET;

```

4. Data Transformation Commands

```

-- Populate dim_users\INSERT OVERWRITE TABLE dim_users
SELECT DISTINCT user_id, region, device
FROM external_user_activity;

-- Populate dim_media
INSERT OVERWRITE TABLE dim_media
SELECT DISTINCT *
FROM external_media_metadata;

-- Populate fact_user_actions
SET hive.exec.dynamic.partition.mode=nonstrict;
SET hive.exec.dynamic.partition=true;

INSERT OVERWRITE TABLE fact_user_actions PARTITION (year, month, day)
SELECT user_id, media_id, action, session_id, event_timestamp, year, month, day
FROM external_user_activity;

```

5. Sample Queries and Execution results:

Query 1: Count of Unique Active Users per Day

```

SELECT year, month, day, COUNT(DISTINCT user_id) AS active_users
FROM user_activity_logs
GROUP BY year, month, day
ORDER BY year DESC, month DESC, day DESC;

```

| year | month | day | active_users |
|------|-------|-----|--------------|
| 2025 | 2 | 7 | 25 |
| 2025 | 2 | 6 | 24 |
| 2025 | 2 | 5 | 27 |
| 2025 | 2 | 4 | 23 |
| 2025 | 2 | 3 | 21 |
| 2025 | 2 | 2 | 24 |
| 2025 | 2 | 1 | 24 |
| 2024 | 3 | 25 | 1 |
| 2024 | 3 | 20 | 1 |
| 2024 | 3 | 12 | 1 |
| 2024 | 3 | 10 | 21 |
| 2024 | 3 | 1 | 1 |
| 2024 | 2 | 5 | 1 |
| 2024 | 2 | 3 | 1 |
| 2024 | 2 | 1 | 1 |
| 2024 | 1 | 16 | 1 |
| 2024 | 1 | 15 | 1 |

```

17 rows selected (37.086 seconds)
0: jdbc:hive2://localhost:10000>

```

Query 2: Top Played Content

```

SELECT content_id, COUNT(*) AS play_count
FROM user_activity_logs
WHERE action = 'play'
GROUP BY content_id
ORDER BY play_count DESC
LIMIT 5;

```

| content_id | play_count |
|------------|------------|
| F222 | 1 |
| F106 | 1 |
| D204 | 1 |
| T236 | 1 |
| A101 | 1 |

```

5 rows selected (36.43 seconds)
0: jdbc:hive2://localhost:10000>

```

Query 3: Dimension Table (dim_content)

```

CREATE TABLE dim_content (
    content_id STRING,
    title STRING,
    category STRING,
    length INT,
    artist STRING
)

```

STORED AS PARQUET;

```

0: jdbc:hive2://localhost:10000> CREATE TABLE dim_content (
    content_id STRING,
    title STRING,
    category STRING,
    length INT,
    artist STRING
)
STORED AS PARQUET;
INFO : Compiling command(queryId=hdoop_20250310161249_c05e244f-52cc-4ef5-9292-5050eaeab4d6b): CREATE TABLE dim_content (
content_id STRING,
title STRING,
category STRING,
length INT,
artist STRING
)
STORED AS PARQUET
INFO : Semantic Analysis Completed (retrial = false)
INFO : Created Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hdoop_20250310161249_c05e244f-52cc-4ef5-9292-5050eaeab4d6b); Time taken: 0.01 seconds
INFO : Currency node is disabled, not creating a lock manager
INFO : Executing command(queryId=hdoop_20250310161249_c05e244f-52cc-4ef5-9292-5050eaeab4d6b): CREATE TABLE dim_content (
content_id STRING,
title STRING,
category STRING,
length INT,
artist STRING
)
STORED AS PARQUET
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hdoop_20250310161249_c05e244f-52cc-4ef5-9292-5050eaeab4d6b); Time taken: 0.058 seconds
0 rows affected (0.073 seconds)
0: jdbc:hive2://localhost:10000>

```

6. Design consideration and performance optimization:

a) Star schemas Design:

It optimizes the query performance where **user_activity_logs** serve as a fact table storing detailed user interactions. **Dim_content** is a dimensions table which stores metadata which is stored in a parquet format which improves compression and read efficiency.

b) Data Storage format:

There were 2 ways in which we stored the data , Fact table and external raw table which helps in efficient analytics and simple ingestion and preprocessing.

c) Query Execution:

Sorting by usage of queries used by **GROUP BY** which we prune the unnecessary data, and we use only year, month and day.

7. Execution Time Analysis:

| Stage | Execution Time |
|----------------------------------|-------------------------|
| Data Ingestion from HDFS | 10-15 seconds |
| Raw Table Creation | 15 seconds |
| Transforming Raw Data to Parquet | 25-40 seconds per table |
| Total Execution Time | 1 minute |

8. Conclusion:

The **Hadoop** and **Hive-based** pipeline efficiently processes and analyzes large datasets using partitioning, Parquet storage, and optimized queries. With a total execution time of 1 minute, the system ensures fast data retrieval and scalability. These design choices enhance performance, and support seamless data-driven decision-making.